

Evaluating Safety and Coordination in Context-Aware Speed and Separation Monitoring for Human–Robot Ceiling-Panel Installation

Zenan Wu, Luke Siniakov, and Michael Magila

Abstract—Human–robot collaboration in ceiling-panel installation requires coordinated movement around a shared payload. Fixed protective zones are straightforward to implement but do not distinguish approaching, working and retreating behaviour. This paper presents a context-aware speed and separation monitoring prototype that combines geometric constraints, task-phase recognition and short-horizon intrusion prediction. A three-state Gaussian-mixture hidden Markov model supplies task context, while a deterministic command ceiling prevents that context from increasing the speed permitted by the reactive envelope in SSM mode. Development validation across three operators achieved causal phase accuracy of 0.794 and balanced accuracy of 0.725. An exploratory evaluation analysed fourteen telemetry trials and 31 block-questionnaire submissions. Responses generally indicated high stopping confidence and low perceived workload; eleven matched intake/end records showed mean safety ratings of 3.82 before and 4.45 after the session. Telemetry identified sensitivity to geometric reference selection and variation in the execution of cued movements. The results demonstrate the integration of context-aware control and multimodal measurement for the panel-installation task, while incomplete participant linkage and unvalidated protective-response measurements limit controller-effect conclusions.

Index Terms—Human–robot collaboration, speed and separation monitoring, hidden Markov model, construction robotics, motion capture, human factors.

I. Introduction and Background

OVERHEAD panel installation combines handling, alignment and fastening in a workspace shared by a worker and a moving payload. A robot can support the panel while the worker supplies judgement and dexterity. However, the worker’s position changes during the task, the payload enlarges the occupied volume, and unnecessary stops can disrupt coordination.

Speed and separation monitoring (SSM) depends on protective separation and timely protective action [1]–[3]. A fixed threshold is transparent but does not distinguish approaching, stationary and retreating motion. Context may improve coordination, provided recognition errors cannot relax the deterministic command bound. This project therefore separates task-phase recognition, kinematic prediction and geometric constraints.

The research question is whether task context can improve coordination beyond fixed-zone and reactive policies while remaining bounded by the same geometric

constraint. The intended comparison considers intrusion response, unnecessary interruption, task efficiency and participant experience. The present exploratory analysis addresses three supporting questions: how accurately the model recognises task phase under causal operation; how geometry and cued movement affect the interpretation of recorded control behaviour; and what the questionnaire responses show about the experience of successive task blocks. These questions characterise the prototype and the evidence needed to evaluate the original comparative aim.

The work contributes a physical ceiling-panel collaboration prototype that separates task-phase recognition from kinematic intrusion prediction under an explicit command ceiling; a descriptive analysis of geometric inputs and robot motion during the task; and item-level human-experience results with a complete-response sensitivity analysis. Its contribution is application-specific integration and evaluation. It does not introduce adaptive SSM itself or establish a controller-specific safety or efficiency advantage.

II. Related Work

A. Construction HRC

Robotic construction research has increasingly moved from isolated automation toward systems in which human expertise and robot strength are combined [4], [5]. Reviews of construction HRC identify safety, perception, communication and task allocation as continuing barriers to deployment [6], [7]. Panel installation is an appropriate case study because the person and robot must coordinate around a shared object and because excessive conservative stopping directly affects task fluency.

B. Speed and Separation Monitoring

SSM requires more than drawing a circle around a robot. The required separation depends on human and robot motion during sensing, controller reaction and robot stopping, together with intrusion and position-uncertainty allowances [2], [3]. Byner et al. [8] derive robot velocity limits using separation and motion direction, and evaluate cycle performance in a physical machine-tending task. Karagiannis et al. [9] switch safety-zone configurations according to robot position, using an overhead safety

camera, safety PLC and robot controller in an automotive-inspired cell. These studies establish that adaptation and productivity evaluation already belong to SSM research. The controller developed here is standards-informed but is not a certified SSM function: its simplified equation and software output must not be confused with validation of the installed robot, tool, payload, sensing uncertainty or safety-rated stop chain.

Construction-specific work further narrows the contribution. Xiao Lin et al. [10] integrate worker, robot and material geometry in a digital twin with dynamic protective-distance calculation, evaluating static and dynamic conflict scenarios. Their treatment of manipulated materials is particularly relevant to panel installation: a proxy around the robot alone does not describe the entire occupied volume. The present head/segment-origin proxies therefore support control diagnostics but cannot establish equivalent geometric coverage. Table I distinguishes the evaluated capabilities of these approaches from those of this prototype.

C. Human Activity Recognition and Collaboration

Hidden Markov models (HMMs) provide an interpretable probabilistic representation of temporal state. They describe a hidden state sequence through a transition matrix and relate each state to observable measurements through an emission model [11]. HMMs have been used in wearable activity recognition and human-robot task modelling [12], [13]. However, task phase and hazard are different variables: a person can trip while approaching, working or retreating. This project therefore uses the HMM only for task context and uses an independent geometric predictor for the operational rapid-intrusion event.

Semantic states such as “working” or “retreating” describe task activity, whereas a hazardous situation concerns exposure to harm. A rapid approach may occur within any phase, so treating hazard as a mutually exclusive phase would confound activity with event recognition. Separating these variables keeps the geometric command bound testable independently of the HMM’s interpretation.

A closely related construction study by Xiao Lin et al. [14] combines a semantic workflow model, policies learned from demonstration and local safety control in a virtual-reality wall-panel task with simulated sensors. Consequently, combining task semantics with safety constraints is not unique to this project. The present study instead examines a physical ceiling-panel prototype with measured robot telemetry and questionnaire responses. Those measurements address implementation and experience; they do not establish superiority over the semantic method or reproduce its policy-learning evaluation.

Objective collaboration fluency is commonly evaluated through temporal measures such as human idle time, robot idle time, concurrent activity and functional delay [15]. Perceived workload and trust can add explanatory

evidence, but subjective reports cannot establish physical safety or replace measured response [16], [17].

The resulting evaluation focus is the connection between task recognition, measured control inputs and perceived experience in a physical panel task. Continuous speed adaptation, workflow intelligence and physical safeguarding are distinct capabilities; evidence for one cannot establish the others.

III. System Design

A. Robot, Sensing and Geometry Versions

The laboratory prototype uses a UR10 robot and a lightweight ceiling-panel surrogate. OptiTrack supplies a head rigid-body position; the Xsens path records 23 segment positions and orientations. Robot telemetry supplies TCP pose and joint velocities. The nominal dashboard loop is 60 Hz, although actual intervals vary. The trial streams, manifests and event journals retain the observed timing.

For a tracked point $\mathbf{p}$ and a nearest point $\mathbf{q}$ on a vertical column proxy, the head feature is

$$d = \|\mathbf{q} - \mathbf{p}\|_2, \quad v_{\text{proj}} = \mathbf{v}^T \mathbf{u}, \quad (1)$$

where $\mathbf{u}$ points toward the proxy and positive projected velocity means closing motion. A recent-sample slope estimates velocity; heading alignment is velocity alignment, not torso orientation.

The analysed trials contain two geometry implementations. P17 and P22 use head features that are not arithmetically consistent with a same-row live-TCP column reconstruction. P24 records a live TCP reference for the head feature and an anchored set of segment origins for the controller’s body-distance input. The phase model still uses head-derived features. These versions are reported separately: neither a head point nor segment origins establish minimum clearance to every limb, robot link, tool or panel surface.

B. Task Context and Intrusion Prediction

The three hidden phases are approaching, working and retreating. The observation vector is $\mathbf{x}_t = [d_t, v_{\text{proj},t}, \|\mathbf{v}_t\|, h_t]^T$, where h_t is velocity-heading alignment. Each state’s emission density is a three-component diagonal Gaussian mixture:

$$p(\mathbf{x}_t \mid z_t = i) = \sum_{k=1}^3 w_{ik} \mathcal{N}(\mathbf{x}_t; \boldsymbol{\mu}_{ik}, \text{diag}(\boldsymbol{\sigma}_{ik}^2)). \quad (2)$$

The online forward filter uses present and past observations; offline Viterbi decoding also uses subsequent observations. Their validation values are therefore reported separately. Long labelled phases and the configured transition transformation produce highly persistent states, a potential limitation during short movements.

The intrusion channel operates independently of the HMM. It solves

$$d - r_{\text{red}} - v_{\text{proj}}\tau - \frac{1}{2}a_{\text{proj}}\tau^2 = 0 \quad (3)$$

TABLE I

Selected related methods and their relationship to the present evaluation. Entries describe reported scope, not a ranking or exhaustive literature survey.

Study	Method and evaluation setting	Implication for this study
Byner et al. [8]	Continuous velocity limits from separation and motion direction; physical machine-tending cycle evaluation.	Reactive speed adaptation is an established comparator; phase recognition must demonstrate incremental value.
Karagiannis et al. [9]	Robot-position-dependent zone switching; overhead safety camera and safety PLC in an automotive-inspired cell.	A software command bound is distinct from a safety architecture implemented with safety components.
X. Lin et al. [10]	Dynamic protective distance with worker, robot and material geometry; construction conflict scenarios.	Panel and body geometry require validation beyond point-proxy arithmetic consistency.
X. Lin et al. [14]	Learned semantic workflow and local safety control; VR wall-panel installation with simulated sensors.	Task-aware safety control already exists; physical telemetry and reported experience define the present evaluation focus.
Present study	Causal HMM phase context and kinematic prediction beneath a command ceiling; physical ceiling-panel prototype.	Exploratory geometry, motion and questionnaire results; comparative protective performance remains unestablished.




Fig. 1. Prototype architecture. In SSM mode, learned outputs can add caution but cannot raise the command above the deterministic envelope. Logging spans sensing, decisions, commands, and measured robot response. The software bound is not a certified safety function.

for a non-negative time within a 0.5-s horizon. The configured closing-speed gate is 0.6 m/s and projected acceleration is limited to $\pm 4 \text{ m/s}^2$. This predicts approach to a software boundary; it does not recognise arbitrary falls or infer intent.

C. Controller Policies and Command Path

The fixed controller maps constant distance bands to full speed, reduced speed and a zero request. The reactive controller uses closing speed in a simplified envelope. The predictive controller adds phase caution and intrusion prediction to that envelope. Its separation parameterisation is

$$S = \max(0, v_{\text{proj}})T + C + S_a, \quad (4)$$

with configured $T = 0.4 \text{ s}$, $C = 0.2 \text{ m}$ and $S_a = 0.1 \text{ m}$. The ramp above S is 0.3 m. The fixed inner radius is 0.94 m; the fixed yellow radius is 1.6 times that value. These are research settings, not measured installed-system stopping or uncertainty bounds.

In SSM mode, the software constraint is

$$u_{\text{cmd}} = \min(u_{\text{env}}, u_{\text{context}}), \quad 0 \leq u_{\text{cmd}} \leq 1, \quad (5)$$

with a fixed inner-boundary override and conservative handling of unavailable tracking. Separate monitored-stop and hand-guide modes exist in the software; the SSM inequality is not a validation of those modes or of permitted

contact. The live command path uses the UR speed slider and Dashboard pause. A software acknowledgement does not establish braking performance, and Dashboard pause is not a safety-rated protective stop. The prototype must therefore be distinguished from a certified installed SSM function.

IV. Model Development and Validation

The phase classifier was developed using 32 accepted labelled trials from P03 (eight), P04 (twelve) and P05 (twelve), containing 90,487 frames. Class counts are 17,267 approaching, 53,609 working and 19,611 retreating. The reported validation holds out one development operator at a time. The values reported here are the stored cross-validation results for the model used in the experiment.

Offline accuracy is 0.821 and causal accuracy is 0.794 (Table II). Balanced accuracy is lower, particularly for the causal filter (0.725), and the working class has the highest offline recall. The small development pool, temporal correlation and use of these folds during development restrict generalisation. These values quantify phase-recognition diagnostics, not intrusion sensitivity or physical safety.

V. Experimental Method

A. Task and Experimental Design

The protocol specifies three blocks labelled A, B and C, with fixed, reactive and predictive conditions assigned

TABLE II

Leave-one-operator-out phase-recognition validation on the development set. Viterbi decoding uses the full sequence; the causal filter uses present and past observations.

Measure	Value
Development participants	3
Accepted trials	32
Labelled frames	90487
Offline Viterbi accuracy	0.821
Offline balanced accuracy	0.757
Causal-filter accuracy	0.794
Causal-filter balanced accuracy	0.725
Approaching recall	0.693
Working recall	0.923
Retreating recall	0.654
Offline fold P03 accuracy	0.812
Offline fold P04 accuracy	0.817
Offline fold P05 accuracy	0.832

through six rotated orders. Each block schedules one clean, one distractor and one rapid-movement trial. The ceiling-panel sequence includes approach and loading, retreat before lifting, work near the raised panel, and retreat before lowering and unloading. The participant pack places the scripted cue during the lift; execution was operator-paced. Planned labels and dashboard cue windows are not independent verification of actual movement onset.

A descriptive analysis was used because complete within-person controller sequences and consistent participant linkage were unavailable. The analysis included 14 telemetry trials associated with three study IDs and 31 block questionnaire submissions associated with eleven IDs. Nine IDs had all three block questionnaires and two lacked the final block response. Matched intake and end responses were available for all eleven follow-up IDs. These counts describe analysed records; limitations of individual-level linkage are discussed in Section VII-D.

Telemetry and questionnaire outcomes were analysed separately because their coverage differs. Recorded samples and original form responses were used directly; no missing values were imputed. Simulation results were not included in the experimental analysis.

B. Telemetry Processing

A trial entered the telemetry analysis when its sample stream and event journal were readable and its sample count matched the trial metadata. Fourteen trials met these criteria. One additional trial record lacked its sample stream and was excluded. The resulting dataset contains 65,697 samples. This criterion establishes the availability of automatic recordings; it does not imply that native motion-capture files or synchronized video exist for every trial.

For timestamps t_i , time summaries use left-held values only when $0 < t_{i+1} - t_i \leq 0.25$ s. Longer intervals are reported separately and excluded from both numerator and denominator. The final sample has no following exposure interval. Fresh motion telemetry requires an available packet, six finite joint velocities and a source

age between zero and 0.25 s. A maximum absolute joint velocity exceeding 0.005 rad/s is a diagnostic motion flag; counts at 0.001 and 0.01 rad/s are retained as threshold sensitivity checks. This flag does not define verified stopping or task completion.

Head-feature consistency is assessed by recomputing point-to-column distance from the same row's head position and live TCP. This is an arithmetic check, not an external calibration. Cue-window extrema describe recorded controller inputs; extrema need not occur at the same instant. Zero-command time includes intentional holds, so it is not classified as false interruption.

C. Questionnaire Inclusion and Scoring

Questionnaire analysis excluded pilot, link-test and qualification submissions. There were 18 intake records, 31 block responses and 11 end responses after those exclusions. Block counts were A: 11, B: 11 and C: 9, with no duplicate ID/block keys. Analyses used the original responses rather than spreadsheet-derived composite scores.

The block questionnaire uses five-point ratings. Three semantic differentials run from anxious to relaxed, agitated to calm, and quiescent to surprised. Four agreement items concern stopping confidence, constant monitoring, unnecessary stopping and speed legibility. Three adapted items concern mental demand, physical demand and frustration. The last three are not a full NASA-TLX score [16], and individual confidence/monitoring items are not the validated trust scale of Jian et al. [17]. All items are retained in their original direction; no composite combines the opposing surprised item with relaxed and calm scores.

Block results are presented as raw response counts and median [first, third quartile], using linear quantile interpolation. A sensitivity summary uses only the nine complete A–C response sets. Within these sets, counts of increases, ties and decreases between A and C retain the pairing and original item direction; they describe sequence-related changes without interpreting higher as uniformly better. Results are grouped by recorded block label rather than controller, because assignment records alone do not establish actual exposure for every questionnaire. Intake/end comparisons match study IDs and retain the change from expected to experienced safety and trust. Both forms use 1 = not at all and 5 = completely.

The analysis is descriptive and retrospective. Neither the independent sample size nor complete controller exposure can be established, so no significance tests, confidence intervals or causal effect estimates are reported. Missing observations are not assumed to be random.

VI. Results

A. Telemetry Characteristics

The telemetry dataset contains five trials under P17, eight under P22 and one under P24. P17 and P22 also have block and end questionnaires; P24 has an intake record. Table III summarises the recorded controller, scheduled

TABLE III

Analysed telemetry trials. Span is capture duration, not isolated task time. Zero-request time includes intentional holds; it is not a false-stop rate.

Code/trial	Controller	Planned event	Rows	Span (s)	Gaps (s)	Zero request (%)
P17 T05	Fixed	Distractor	5,186	103.2	2.45	30.1
P17 T06	Fixed	Clean	4,550	91.7	2.55	30.8
P17 T07	Reactive	Clean	5,061	101.1	2.88	41.3
P17 T08	Reactive	Distractor	4,815	97.0	3.06	34.2
P17 T09	Reactive	Rapid Cue	4,426	92.1	2.50	32.8
P22 T02	Reactive	Distractor	5,415	108.3	2.31	34.4
P22 T03	Reactive	Rapid Cue	4,668	93.6	2.28	29.0
P22 T05	Predictive	Rapid Cue	5,315	109.0	3.05	34.9
P22 T06	Predictive	Clean	4,400	90.0	2.62	27.0
P22 T07	Predictive	Distractor	4,950	101.1	2.75	25.9
P22 T08	Fixed	Rapid Cue	4,745	96.0	3.13	31.4
P22 T09	Fixed	Distractor	4,772	95.6	2.58	24.9
P22 T10	Fixed	Clean	4,761	95.1	2.59	25.3
P24 T01	Predictive	Rapid Cue	2,633	114.2	4.45	39.9

Intervals over 250 ms are excluded from time integrals. Zero-request percentage uses the remaining positive intervals as its denominator. P17 and P22 use the earlier head geometry; P24 uses anchored segment origins for control.

event, capture span and zero-command proportion for each analysed trial.

The 14 analysed trials contain 65,697 samples and 23.13 minutes of summed capture spans (Table III). 57 intervals exceed 250 ms, totalling 39.20 s or 2.82% of those spans; 1348.69 s remain for time integrals. Forty rows carry a stale-tracking flag. Controller decisions exist in 65,684 rows. One hundred and ten decisions indicate unavailable tracking and invoke the conservative controller state. These counts characterise the sampled command stream; physical response is considered separately.

Across ten non-clean cue windows, fresh joint telemetry exceeds the diagnostic motion threshold for 236.31 s of observed intervals. The measurements confirm robot motion during labelled event windows, although the operator cue does not establish the onset or severity of physical intrusion. Among 13,048 fresh telemetry rows in these windows, motion counts are 11,509, 11,355 and 11,319 at 0.001, 0.005 and 0.01 rad/s, respectively. The presence of motion is therefore not dependent on the selected threshold alone. These are correlated sample counts, not independent events or tests of stopping performance.

B. Geometric Inputs and Cued Movements

The P17/P22 trial-level 95th-percentile discrepancies between recorded head distance and same-row live-TCP recomputation range from 0.413 to 0.534 m (Fig. 2a). P24's corresponding discrepancy is below 10^{-12} m, showing arithmetic agreement for its head reference. This improvement does not validate the anchored segment-origin geometry or its physical uncertainty.

Recorded cue labels also differ from their intended kinematics. P17 T09 is labelled rapid intrusion, yet its minimum controller proxy is 2.034 m and maximum closing input is 0.388 m/s, below the configured gate. Conversely, the P17 T08 distractor window reaches a 0.496-m proxy and a 1.494-m/s maximum closing input. P22 T05 reaches a 0.252-m proxy during a rapid cue. These values identify records requiring interpretive caution; they

are not validated human-robot clearances or incident classifications.

Figure 3 presents the recorded time series for P24 T01. Its rapid-cue window reaches a 0.913-m controller proxy and a 5.302-m/s maximum closing input. The large peak and changed body geometry make independent motion validation particularly important. Command and joint-speed traces are presented together, but no delay measured from the operator's cue is called protective response time.

C. Reported Experience

Block responses lean toward relaxed/calm, confidence in stopping, and low demand and frustration (Table IV, Fig. 4). Median stopping confidence is five in each block, while median unnecessary-stopping agreement is three in A and two in B/C. The latter is a perception item, not an objective false-stop count. Pooled distributions contain repeated responses from the same codes and are not a sample of 31 people.

Restricting the analysis to nine complete response sets preserves 29 of the 30 block-item medians. The exception is constant monitoring in B, whose median changes from two to one. The paired A-C counts nevertheless reveal variation hidden by marginal medians (Table VI). Five relaxed ratings increase and four remain unchanged; none decrease. Unnecessary-stopping ratings increase twice, decrease twice and remain unchanged five times, despite the lower C median. Surprise decreases for three codes, increases for two and is unchanged for four. Thus lower block medians do not establish a common within-code improvement.

Across eleven matched records, mean safety ratings were 3.82 before and 4.45 after the session. Eight ratings increased, one was unchanged and two decreased. Mean trust ratings were 4.00 before and 4.18 after; five increased, three were unchanged and three decreased (Table V). The expected-to-experienced wording and incomplete individual linkage limit interpretation of these differences.




Fig. 2. Measurement diagnostics. (a) Head features differ from the same-row live-TCP column calculation. (b) Independent extrema within cue windows; dashed lines show the configured 0.94-m boundary and 0.6-m/s closing gate. Colours identify fixed (blue), reactive (orange) and predictive (green) records. Geometry versions differ and extrema are not necessarily simultaneous.




Fig. 3. P24 T01: recorded body-proxy distance, closing input, command fraction and maximum absolute fresh joint velocity. Shading marks the dashboard event window. Dashed lines are configuration thresholds. Gaps over 250 ms break the lines. The plot is an observed trace, not a counterfactual controller comparison.

In end responses, C is selected as safest seven times, compared with two each for A and B. Full-shift choices are C: six, B: three and A: two. Least-safe choices are distributed across all labels (A: three; B: four; C: four). Since label-to-controller exposure is not fully established, C is not interpreted as the predictive controller. Four end responses report noticing changes and seven do not. The four comments refer to stopping/slowing sensitivity or ascent speed; they are descriptive observations, not a formally coded thematic analysis. Median reported suit restriction is two (interquartile range 1.5–3). No risk

or discomfort comments were entered in the eleven end forms; these optional responses are not an independently observed incident measure.

VII. Discussion

A. Context-Aware Control in the Panel Task

The prototype integrates task context and kinematic prediction while preserving a deterministic software ceiling on the commanded speed. This architecture separates two questions: whether a movement belongs to a normal task phase, and whether the measured trajectory approaches

TABLE IV

Block questionnaire items: median [first, third quartile]. A, B and C are reported block labels, not pooled controller conditions. All ratings use 1–5.

Item	A ($n = 11$)	B ($n = 11$)	C ($n = 9$)
Anxious to relaxed	4 [3.5, 4.5]	5 [4.0, 5.0]	5 [4.0, 5.0]
Agitated to calm	4 [4.0, 5.0]	5 [4.0, 5.0]	5 [4.0, 5.0]
Quiescent to surprised	3 [1.0, 3.0]	3 [1.0, 4.0]	1 [1.0, 3.0]
Confident robot would stop	5 [4.0, 5.0]	5 [4.0, 5.0]	5 [4.0, 5.0]
Had to watch constantly	2 [1.0, 3.0]	2 [1.0, 3.0]	1 [1.0, 3.0]
Stopped more than needed	3 [2.0, 4.0]	2 [1.0, 3.5]	2 [2.0, 3.0]
Speed changes made sense	5 [3.5, 5.0]	5 [4.0, 5.0]	4 [4.0, 5.0]
Mental demand	1 [1.0, 3.0]	1 [1.0, 2.0]	1 [1.0, 2.0]
Physical demand	1 [1.0, 3.0]	1 [1.0, 2.5]	1 [1.0, 3.0]
Frustration	1 [1.0, 3.0]	1 [1.0, 2.0]	1 [1.0, 1.0]

Higher means the right-hand semantic anchor for the first three items, stronger agreement for the next four, and greater demand/frustration for the final three. Responses are clustered within 11 study IDs; individual-level linkage is incomplete. No trust, safety or workload composite is used.




Fig. 4. Original item distributions across 31 block submissions. Each bar sums to 31; embedded values are response counts. Higher scores mean more relaxed/calm, more surprised, stronger agreement, or greater demand/frustration according to the item. Pooling describes the recorded responses and does not rank controllers.

TABLE V

Paired global ratings for 11 codes with intake and end responses. Five-point items; descriptive code-level changes only.

Item	Before	After	Δ	Up/same/down
Felt safe	3.82	4.45	0.64	8/1/2
Trust in distance	4.00	4.18	0.18	5/3/3

Before/after wording changes from expected to experienced safety/trust. Means summarize coded responses; neither causation nor an interval-scale measurement model is assumed.

the configured boundary. An HMM misclassification need not authorise a higher speed than the geometric envelope, although errors in geometry or timing still affect that envelope.

The development results indicate that phase recognition is strongest during the longer working phase. The lower causal balanced accuracy compared with overall accuracy shows why phase accuracy alone is an insufficient measure

TABLE VI

Complete response sets: block medians and within-code changes ($n = 9$ codes).

Item	A/B/C	A to C: up/same/down
Anxious to relaxed	4/5/5	5/4/0
Agitated to calm	4/5/5	3/6/0
Quiescent to surprised	3/3/1	2/4/3
Confident robot would stop	5/5/5	3/4/2
Had to watch constantly	2/1/1	1/5/3
Stopped more than needed	3/2/2	2/5/2
Speed changes made sense	5/5/4	1/5/3
Mental demand	1/1/1	2/6/1
Physical demand	1/1/1	1/6/2
Frustration	1/1/1	1/5/3

Counts refer to the raw rating direction, not uniformly to improvement. A–C changes describe block sequence, not a controller effect. The same nine codes supply every row.

of practical benefit. Additional caution from a persistent state estimate could increase interruption, particularly

around short approach/retreat transitions. The reactive reference is therefore an important comparator when assessing the added value of learned context.

Equation 5 also constrains the expected benefit. For identical SSM inputs, parameters and zone state, the predictive policy cannot request a higher speed than the reactive policy: prediction and phase caution can only preserve or reduce the request. Its incremental value would therefore need to arise from the timing of intervention or changes in subsequent human–robot coordination, rather than a more permissive instantaneous speed bound. Such a closed-loop benefit cannot be inferred from phase accuracy or from applying alternative policies to the same recorded trajectory.

B. Objective Behaviour and Perceived Experience

The telemetry illustrates changes in the distance input, speed request and measured joint motion throughout the panel sequence. The difference between geometric reference configurations materially affects how these signals can be interpreted. Similarly, variation between scheduled rapid movements and recorded closing inputs shows that event labels cannot substitute for measured exposure. These findings are consistent with the need to validate the sensing-to-stopping chain when implementing SSM [3].

Questionnaire responses generally indicate confidence in the robot stopping, low mental and physical demand, and limited frustration. Mean experienced safety exceeds mean expected safety in the matched records, whereas the trust difference is smaller. These observations suggest a positive reported experience of the laboratory task. However, novelty, supervised operation, repeated exposure and expectations could all contribute to the pattern. The change cannot be attributed specifically to the predictive controller.

The complete-response analysis makes the broad descriptive pattern less dependent on the two missing final-block forms, but does not remove bias from missingness or resolve individual linkage. Ceiling ratings for stopping confidence and floor ratings for demand/frustration also leave limited room to distinguish conditions. The paired stopping and surprise responses show why aggregate medians should be read alongside within-code changes. Increased familiarity is a plausible explanation for greater relaxation across blocks, but was not isolated by this analysis.

The end-of-session preference for block C does not provide a controller ranking: block labels rotate across controller assignments, and complete exposure linkage is unavailable. Likewise, subjective agreement that the robot stopped more than needed is informative about perceived interruption but is not an objective false-stop rate. Subjective and physical measures address different aspects of collaboration and should be interpreted together.

C. Safety and Efficiency Interpretation

Dynamic SSM research motivates evaluation of availability and cycle performance [8], [9]. The present study

demonstrates the implementation needed to investigate those outcomes but does not establish a safety or efficiency advantage. Zero-command time combines intentional holds and responses to proximity; capture span includes operator actions and cannot be equated with isolated task time. These quantities therefore do not establish false interruption, task efficiency or collaboration fluency [15].

A non-inferiority conclusion would additionally require a justified margin, validated protective geometry and end-to-end response measurements. Those requirements are not met by the current analysis. Absence of a demonstrated difference must not be interpreted as equivalent safety. The contribution is an integrated research system with exploratory observations of behaviour and experience, rather than a validated replacement for an industrial safety function.

D. Limitations

The experimental data are incomplete and derive from a small convenience sample. Session restarts prevented reliable linkage of every study ID to a distinct individual. The numbers of records consequently cannot be treated as independent participant counts. Two questionnaire sets lack block C, and objective and subjective records cover different subsets of the study. Missingness may relate to technical problems or task execution, limiting complete-case interpretation and generalisation.

The geometric configuration changed during collection. Head points and segment origins do not represent complete body, tool and payload surfaces, and their agreement with a geometric calculation does not establish physical calibration. The simplified separation envelope and ordinary robot command interface lack installed-system validation for a certified safety claim. Cued movements, operator pacing, visible motion, learning and fatigue may influence both behaviour and subjective reports. The model-development results are based on three operators and do not establish generalisation to construction workers.

E. Implications for Further Development

Future controlled evaluation should maintain a consistent geometry and validate its uncertainty against the complete moving volume. Synchronized measurements of physical movement, controller output and robot deceleration would allow response-time estimation independently of operator cues. Task-phase estimates could then be evaluated for their incremental benefit over the reactive policy using verified within-person controller exposure. A persistent study identifier across restarted sessions would support that comparison without changing the participant-facing task.

VIII. Conclusion

This study presents a context-aware speed and separation monitoring prototype for human–robot ceiling-panel installation. Its architecture combines measured closing

kinematics, three-state task recognition and intrusion prediction under a deterministic software command ceiling. Causal phase recognition achieved 0.794 accuracy in development validation. An exploratory analysis of fourteen telemetry trials and 31 block responses found generally favourable reported experience and identified geometric-reference and cue-execution issues affecting objective interpretation. Complete-response analysis preserved 29 of 30 block medians, while paired changes showed that lower median interruption ratings did not represent a uniform improvement. The findings support further investigation of context-aware coordination, while a controller-specific safety or efficiency improvement remains unestablished.

Data Availability

Analysis scripts, derived tables and figure-generation code accompany the manuscript. Participant-level data require controlled access. All reported experimental summaries are calculated from recorded observations without imputation or substitution of simulated values.

Acknowledgment

The authors thank Associate Professor Yihai Fang for his primary supervision and academic guidance, Yizhe (Will) Wang for day-to-day technical and laboratory support, and the Construction Robotics Laboratory for providing facilities.

AI Assistance Declaration

AI tools assisted with data-processing code, descriptive analysis, figure production and manuscript drafting and revision.

References

- [1] International Organization for Standardization, “ISO 10218-2:2025 robotics—safety requirements—part 2: Industrial robot applications and robot cells,” International Standard, 2025. [Online]. Available: <https://www.iso.org/standard/73934.html>
- [2] —, “ISO/TS 15066:2016 robots and robotic devices—collaborative robots,” Technical Specification, 2016. [Online]. Available: <https://www.iso.org/standard/62996.html>
- [3] J. A. Marvel and R. Norcross, “Implementing speed and separation monitoring in collaborative robot workcells,” *Robotics and Computer-Integrated Manufacturing*, vol. 44, pp. 144–155, 2017.
- [4] T. Bock, “The future of construction automation: Technological disruption and the upcoming ubiquity of robotics,” *Automation in Construction*, vol. 59, pp. 113–121, 2015.
- [5] M. Zhang, R. Xu, H. Wu, J. Pan, and X. Luo, “Human–robot collaboration for on-site construction,” *Automation in Construction*, vol. 150, p. 104812, 2023.
- [6] H.-H. Wei, Y. Zhang, X. Sun, J. Chen, and S. Li, “Intelligent robots and human–robot collaboration in the construction industry: A review,” *Journal of Intelligent Construction*, vol. 1, no. 1, p. 9180002, 2023.
- [7] P. Lin, N. Zeng, Q. Li, and K. Nübel, “A review of human–robot collaboration safety in construction,” *Systems*, vol. 13, no. 10, p. 856, 2025.
- [8] C. Byner, B. Matthias, and H. Ding, “Dynamic speed and separation monitoring for collaborative robot applications—concepts and performance,” *Robotics and Computer-Integrated Manufacturing*, vol. 58, pp. 239–252, 2019.
- [9] P. Karagiannis, N. Kousi, G. Michalos, K. Dimoulas, K. Mparis, D. Dimosthenopoulos, Ö. Tokçalar, T. Guasch, G. P. Gerio, and S. Makris, “Adaptive speed and separation monitoring based on switching of safety zones for effective human robot collaboration,” *Robotics and Computer-Integrated Manufacturing*, vol. 77, p. 102361, 2022.
- [10] X. Lin, Z. Guo, X. Jin, and H. Guo, “Digital twin-enabled safety monitoring system for seamless worker–robot collaboration in construction,” *Automation in Construction*, vol. 174, p. 106147, 2025.
- [11] L. R. Rabiner, “A tutorial on hidden Markov models and selected applications in speech recognition,” *Proceedings of the IEEE*, vol. 77, no. 2, pp. 257–286, 1989.
- [12] N. Manouchehri and N. Bouguila, “Human activity recognition with an HMM-based generative model,” *Sensors*, vol. 23, no. 3, p. 1390, 2023.
- [13] J. Berg, T. Reckordt, C. Richter, and G. Reinhart, “Action recognition in assembly for human–robot-cooperation using hidden Markov models,” *Procedia CIRP*, vol. 76, pp. 205–210, 2018.
- [14] X. Lin, Z. Guo, H. Guo, and X. Jin, “Safety-aware behavioral decision-making for construction robots: Semantic-driven hybrid framework for worker–robot collaboration,” *Automation in Construction*, vol. 189, p. 107067, 2026.
- [15] G. Hoffman, “Evaluating fluency in human–robot collaboration,” *IEEE Transactions on Human-Machine Systems*, vol. 49, no. 3, pp. 209–218, 2019.
- [16] S. G. Hart and L. E. Staveland, “Development of NASA-TLX (task load index): Results of empirical and theoretical research,” in *Human Mental Workload*, P. A. Hancock and N. Meshkati, Eds. Amsterdam: North-Holland, 1988, pp. 139–183.
- [17] J.-Y. Jian, A. M. Bisantz, and C. G. Drury, “Foundations for an empirically determined scale of trust in automated systems,” *International Journal of Cognitive Ergonomics*, vol. 4, no. 1, pp. 53–71, 2000.